

# tidb压测题目

分值：300 题目描述：使用 sysbench、go-ycsb 和 go-tpc 分别对 TiDB 进行测试并且产出测试报告。测试报告需要包括以下内容：

- 部署环境的机器配置(CPU、内存、磁盘规格型号)
- 拓扑结构(TiDB、TiKV 各部署于哪些节点)
- 调整过后的 TiDB 和 TiKV 配置
- 测试输出结果记录
- 关键指标的监控截图：
  - TiDB Query Summary 中的 qps 与 duration
  - TiKV Details 面板中 Cluster 中各 server 的 CPU 以及 QPS 指标
  - TiKV Details 面板中 grpc 的 qps 以及 duration
- 输出：写出你对该配置与拓扑环境和 workload 下 TiDB 集群负载的分析，提出你认为的 TiDB 的性能的瓶颈所在(能提出大致在哪个模块即可)
- 截止时间：下周二 ( 8.25 ) 24:00:00(逾期提交不给分)

## 部署环境的机器配置(CPU、内存、磁盘规格型号)

### 两台vmwar虚拟机，各8C,16G

```
[root@tbase01 ~]# lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                 8
On-line CPU(s) list:   0-7
Thread(s) per core:    1
Core(s) per socket:    8
Socket(s):              1
NUMA node(s):          1
Vendor ID:              GenuineIntel
CPU family:             6
Model:                 44
Model name:             Intel(R) Xeon(R) CPU           E5645   @ 2.40GHz
Stepping:               2
CPU MHz:                2393.703
BogoMIPS:               4788.00
Hypervisor vendor:      VMware
Virtualization type:    full
L1d cache:              32K
L1i cache:              32K
L2 cache:               256K
L3 cache:               12288K
```

```
NUMA node0 CPU(s):      0-7
[root@tbase01 ~]#
[root@tbase01 ~]# free -h
```

	total	used	free	shared	buff/cache	available
Mem:	15G	345M	13G	628M	2.0G	14G
Swap:	0B	0B	0B			

```
[root@tbase01 ~]#
```

## 各分配了一块150G虚拟硬盘

```
[root@tbase01 ~]# cat /proc/scsi/scsi
Attached devices:
Host: scsi0 Channel: 00 Id: 00 Lun: 00
  Vendor: VMware   Model: Virtual disk   Rev: 2.0
  Type:   Direct-Access               ANSI SCSI revision: 06
Host: scsi3 Channel: 00 Id: 00 Lun: 00
  Vendor: NECVMwar Model: VMware SATA CD00 Rev: 1.00
  Type:   CD-ROM                      ANSI SCSI revision: 05
[root@tbase01 ~]#
[root@tbase02 ~]# ls SCSI
[0:0:0:0] disk VMware virtual disk 2.0 /dev/sda
[3:0:0:0] cd/dvd NECVMwar VMware SATA CD00 1.00 /dev/sr0
[root@tbase02 ~]# hdparm -i /dev/sda
/dev/sda:
SG_IO: bad/missing sense data, sb[]: 70 00 05 00 00 00 00 0a 00 00 00 00 20 00
00 c0 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
HDIO_GET_IDENTITY failed: Invalid argument
[root@tbase02 ~]#
```

## 测试硬盘

### 随机同步写入性能

```
[root@tbase02 ~]# time dd if=/dev/zero of=/tmp/test bs=8k count=51200
oflag=dsync
51200+0 records in
51200+0 records out
419430400 bytes (419 MB) copied, 138.307 s, 3.0 MB/s

real    2m18.315s
user    0m0.106s
sys     0m4.403s
[root@tbase02 ~]#
```

### 批量写入性能

```
[root@tbase02 ~]# time dd if=/dev/zero of=/tmp/test4 bs=8k count=51200
conv=fsync
51200+0 records in
51200+0 records out
419430400 bytes (419 MB) copied, 2.04958 s, 205 MB/s

real    0m2.051s
user    0m0.011s
sys     0m0.541s
[root@tbase02 ~]#
```

## 读出性能

```
[root@tbase02 ~]# free -m
              total        used         free       shared  buff/cache   available
Mem:          15886         1324         4974          1008         9587        13194
Swap:           0           0           0

[root@tbase02 ~]# echo 3 > /proc/sys/vm/drop_caches
[root@tbase02 ~]# free -m
              total        used         free       shared  buff/cache   available
Mem:          15886         1323        13386          1008         1176        13290
Swap:           0           0           0

[root@tbase02 ~]# time dd if=/tmp/test of=/dev/null bs=8k count=51200
oflag=dsync
51200+0 records in
51200+0 records out
419430400 bytes (419 MB) copied, 2.06349 s, 203 MB/s

real    0m2.092s
user    0m0.010s
sys     0m0.342s
[root@tbase02 ~]#
```

## 拓扑结构(TiDB、TiKV 各部署于哪些节点)

```
[root@tbase02 ~]# cat wei_tidb_241and242.yaml
# # Global variables are applied to all deployments and used as the default
value of
# # the deployments if a specific deployment value is missing.
global:
  user: "tidb"
  ssh_port: 22
  deploy_dir: "/tidb-deploy830"
  data_dir: "/tidb-data830"

# # Monitored variables are applied to all the machines.
monitored:
  node_exporter_port: 19100
  blackbox_exporter_port: 19115

server_configs:
  tidb:
    log.slow-threshold: 300
```

```

tikv:
  readpool.storage.use-unified-pool: false
  readpool.coprocessor.use-unified-pool: true
pd:
  replication.enable-placement-rules: true
tiflash:
  logger.level: "info"

pd_servers:
- host: 192.168.1.242

tidb_servers:
- host: 192.168.1.241
- host: 192.168.1.242

tikv_servers:
- host: 192.168.1.241
  port: 30160
  status_port: 30180

- host: 192.168.1.242
  port: 30161
  status_port: 30181

tiflash_servers:
- host: 192.168.1.242

monitoring_servers:
- host: 192.168.1.242

grafana_servers:
- host: 192.168.1.242
[root@tbase02 ~]#
[root@tbase02 ~]# tiup cluster destroy tidb-test
[root@tbase02 ~]# tiup cluster deploy tidb830 v4.0.0 ./wei_tidb_241and242.yaml -
-user root -p
Starting component `cluster`: /root/.tiup/components/cluster/v1.0.9/tiup-cluster
deploy tidb830 v4.0.0 ./wei_tidb_241and242.yaml --user root -p
Please confirm your topology:
tidb cluster: tidb830
tidb version: v4.0.0

```

Type	Host	Ports	OS/Arch	Directories
pd	192.168.1.241	2379/2380	linux/x86_64	/tidb-deploy830/pd-2379,/tidb-data830/pd-2379
tikv	192.168.1.241	30160/30180	linux/x86_64	/tidb-deploy830/tikv-30160,/tidb-data830/tikv-30160
tikv	192.168.1.242	30161/30181	linux/x86_64	/tidb-deploy830/tikv-30161,/tidb-data830/tikv-30161
tidb	192.168.1.241	4000/10080	linux/x86_64	/tidb-deploy830/tidb-4000
tidb	192.168.1.242	4000/10080	linux/x86_64	/tidb-deploy830/tidb-4000
tiflash	192.168.1.242	9000/8123/3930/20170/20292/8234	linux/x86_64	/tidb-deploy830/tiflash-9000,/tidb-data830/tiflash-9000

```
prometheus 192.168.1.242 9090 linux/x86_64 /tidb-
deploy830/prometheus-9090,/tidb-data830/prometheus-9090
grafana 192.168.1.242 3000 linux/x86_64 /tidb-
deploy830/grafana-3000
Attention:
  1. If the topology is not what you expected, check your yaml file.
  2. Please confirm there is no port/directory conflicts in same host.
Do you want to continue? [y/N]: y
```

```
# 这个etcd是什么还不清楚，以后再研究吧。好像是和tbase冲突了
Destroying instance 192.168.1.241
retry error: operation timed out after 1m0s
192.168.1.241 error destroying pd: timed out waiting for port 2379 to be stopped
after 1m0s

Error: failed to destroy pd: 192.168.1.241 error destroying pd: timed out
waiting for port 2379 to be stopped after 1m0s: timed out waiting for port 2379
to be stopped after 1m0s

Verbose debug logs has been written to /root/logs/tiup-cluster-debug-2020-08-22-
10-53-15.log.
Error: run `/root/.tiup/components/cluster/v1.0.9/tiup-cluster`
(wd:/root/.tiup/data/S8MGkCu) failed: exit status 1
[root@tbase02 ~]#
[root@tbase01 bin]# mv /usr/bin/etcd /tmp
[root@tbase01 bin]# ps -ef |grep etcd
etcd      9925      1  2 10:57 ?        00:00:04 /usr/bin/etcd --name=etcd1 --
data-dir=/data/etcd_data --listen-client-
urls=http://0.0.0.0:2379,http://0.0.0.0:4001
root      13468    2142  0 11:00 pts/1    00:00:00 grep --color=auto etcd
[root@tbase01 bin]# kill -9 9925
[root@tbase01 bin]# ps -ef |grep etcd
tbase     13721    2373  0 11:01 ?        00:00:00 bash -c export
ETCDCTL_API=3;etcdctl endpoint health --endpoints=192.168.1.241:2379
tbase     13722    13721  1 11:01 ?        00:00:00 etcdctl endpoint health --
endpoints=192.168.1.241:2379
root      13757    2142  0 11:01 pts/1    00:00:00 grep --color=auto etcd
[root@tbase01 bin]# cp /tmp/etcd /usr/bin/etcd
# 结果：确实是tbase的etcd和tidb-pd冲突，将pd移动到第二台机器上正常。
[root@tbase02 ~]# tiup cluster start tidb830
```

```
[root@tbase02 ~]# tiup cluster display tidb830
Starting component `cluster`: /root/.tiup/components/cluster/v1.0.9/tiup-cluster
display tidb830
tidb cluster: tidb830
tidb version: v4.0.0
```

ID	Role	Host	Ports	OS/Arch	Status	Data Dir	Deploy Dir
192.168.1.242:3000	grafana	192.168.1.242	3000	linux/x86_64	Up	-	/tidb-deploy830/grafana-3000
192.168.1.242:2379	pd	192.168.1.242	2379/2380	linux/x86_64	Up L UI	/tidb-data830/pd-2379	/tidb-deploy830/pd-2379

```

192.168.1.242:9090    prometheus 192.168.1.242 9090
linux/x86_64 up      /tidb-data830/prometheus-9090 /tidb-
deploy830/prometheus-9090
192.168.1.241:4000    tidb        192.168.1.241 4000/10080
linux/x86_64 up      - /tidb-deploy830/tidb-4000
192.168.1.242:4000    tidb        192.168.1.242 4000/10080
linux/x86_64 up      - /tidb-deploy830/tidb-4000
192.168.1.242:9000    tiflash     192.168.1.242 9000/8123/3930/20170/20292/8234
linux/x86_64 up      /tidb-data830/tiflash-9000 /tidb-deploy830/tiflash-
9000
192.168.1.241:30160   tikv        192.168.1.241 30160/30180
linux/x86_64 up      /tidb-data830/tikv-30160 /tidb-deploy830/tikv-30160
192.168.1.242:30161   tikv        192.168.1.242 30161/30181
linux/x86_64 up      /tidb-data830/tikv-30161 /tidb-deploy830/tikv-30161
[root@tbase02 ~]#

```

## 测试环境小结

- 配置环境说明

项目	
操作系统	CentOS Linux release 7.3.1611
TiDB 版本	TiDB-v4.0.0
TiDB & PD & KV	hosts1 : tidb-server , tikv-server ,
TiDB & PD & KV	hosts2 : tidb-server , tikv-server , tipd-server , grafana, prometheus, tiflash
TiDB 默认参数	[log] slow-threshold = 300
TiKV 默认参数	[readpool] [readpool.coprocessor] use-unified-pool = true [readpool.storage] use-unified-pool = false

## 测试输出结果记录

sysbench、go-ycsb 和 go-tpc

### sysbench

这里记录的是point\_select点查询的结果和截图，用A机B机从2个tidb节点同时执行点查：

```

[root@tbase02 wangwei]# sysbench --config-file=tidb.cfg oltp_point_select --
tables=6 --table-size=100000 prepare

```

```

[root@tbase02 wangwei]# sysbench --config-file=tidb.cfg oltp_point_select --
tables=6 --table-size=100000 run
sysbench 1.0.14 (using bundled LuaJIT 2.1.0-beta2)

Running the test with following options:
Number of threads: 8
Report intermediate results every 10 second(s)
Initializing random number generator from current time

Initializing worker threads...

Threads started!

[ 10s ] thds: 8 tps: 12092.49 qps: 12092.49 (r/w/o: 12092.49/0.00/0.00) lat
(ms,95%): 0.87 err/s: 0.00 reconn/s: 0.00
[ 20s ] thds: 8 tps: 11755.09 qps: 11755.09 (r/w/o: 11755.09/0.00/0.00) lat
(ms,95%): 0.90 err/s: 0.00 reconn/s: 0.00
[ 30s ] thds: 8 tps: 11693.12 qps: 11693.12 (r/w/o: 11693.12/0.00/0.00) lat
(ms,95%): 0.92 err/s: 0.00 reconn/s: 0.00
[ 40s ] thds: 8 tps: 11667.91 qps: 11667.91 (r/w/o: 11667.91/0.00/0.00) lat
(ms,95%): 0.92 err/s: 0.00 reconn/s: 0.00
[ 50s ] thds: 8 tps: 11755.78 qps: 11755.78 (r/w/o: 11755.78/0.00/0.00) lat
(ms,95%): 0.92 err/s: 0.00 reconn/s: 0.00
SQL statistics:
  queries performed:
    read: 707485
    write: 0
    other: 0
    total: 707485
  transactions: 707485 (11790.49 per sec.)
  queries: 707485 (11790.49 per sec.)
  ignored errors: 0 (0.00 per sec.)
  reconnects: 0 (0.00 per sec.)

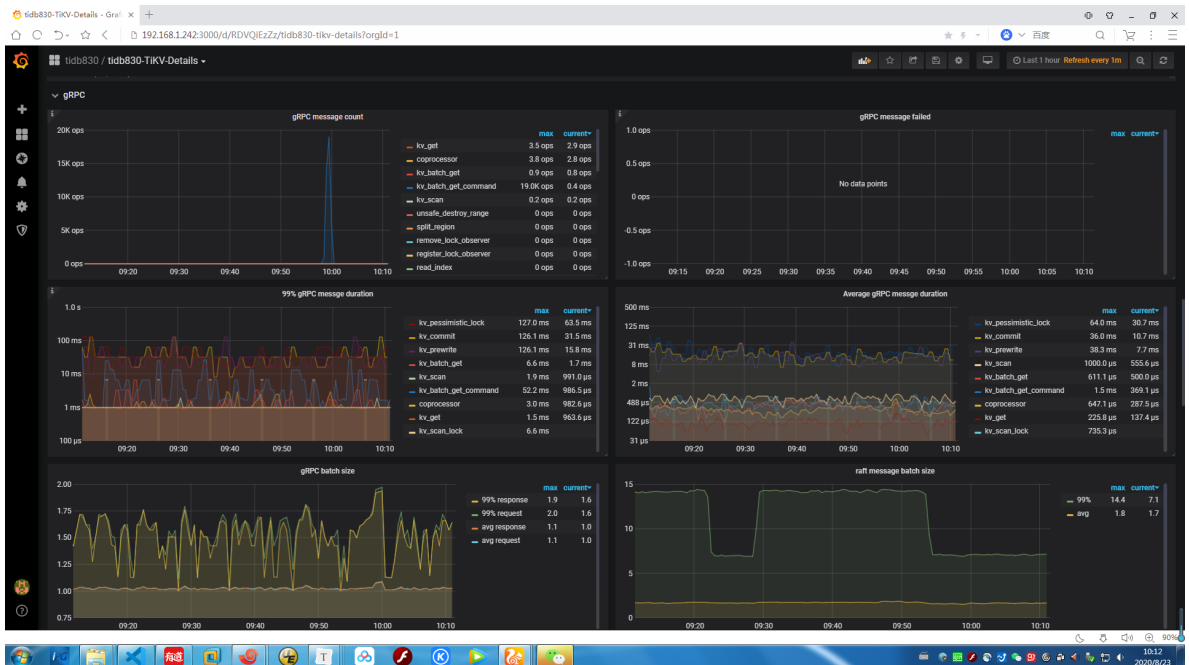

General statistics:
  total time: 60.0026s
  total number of events: 707485

Latency (ms):
  min: 0.35
  avg: 0.68
  max: 23.46
  95th percentile: 0.90
  sum: 478804.31

Threads fairness:
  events (avg/stddev): 88435.6250/24.85
  execution time (avg/stddev): 59.8505/0.00

[root@tbase02 wangwei]#

```





## go-ycsb

```
yum install golang

mkdir /home/gocode/

echo 'export GOPATH=/home/gocode/' >> /root/.bashrc
source /root/.bashrc

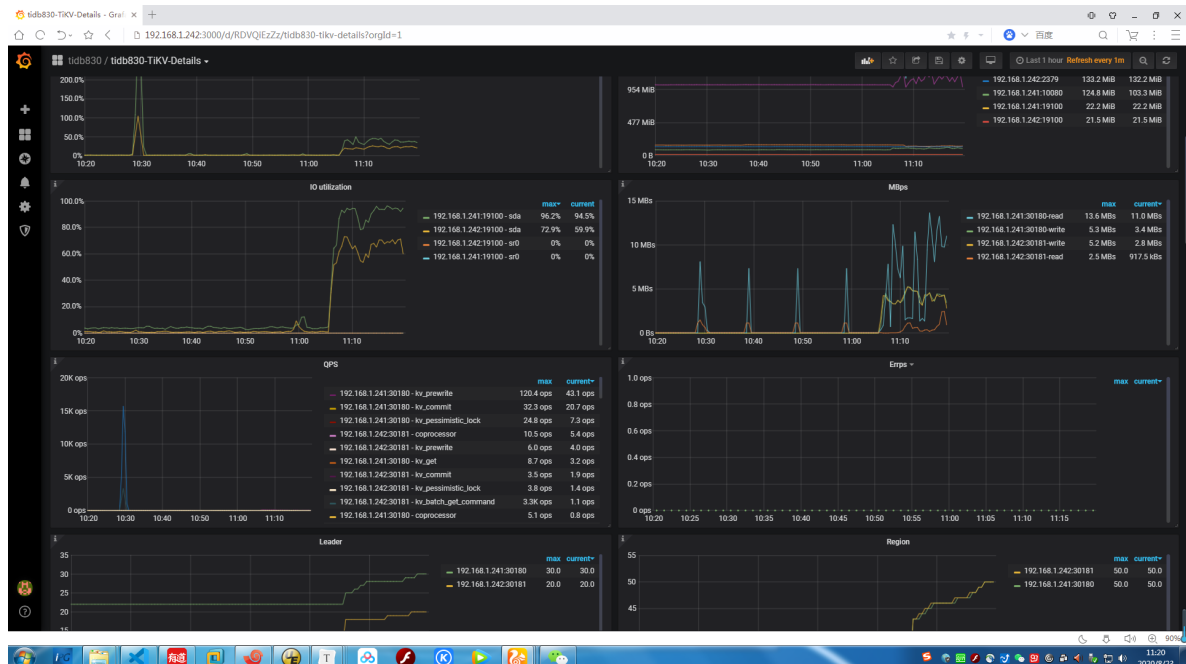
git clone https://github.com/pingcap/go-ycsb.git
$GOPATH/src/github.com/pingcap/go-ycsb
cd $GOPATH/src/github.com/pingcap/go-ycsb
make
```

make安装失败：

```
[root@tbase02 go-ycsb]# make
go build -o bin/go-ycsb cmd/go-ycsb/*
go: github.com/ghodss/yaml@v1.0.1-0.20190212211648-25d852aeb32: invalid pseudo-
version: git fetch --unshallow -f https://github.com/ghodss/yaml in
/home/gocode/pkg/mod/cache/vcs/5c75ad62eb9c289b6ed86c76998b4ab8c8545a841036e879d
703a2bbc5fcfcea: exit status 128:
    fatal: git fetch-pack: expected shallow list
make: *** [build] Error 1
[root@tbase02 go-ycsb]#
```

## go-tpc

加载10个仓库的数据：





## 测试tpcc

做了三次tpmc在400左右。

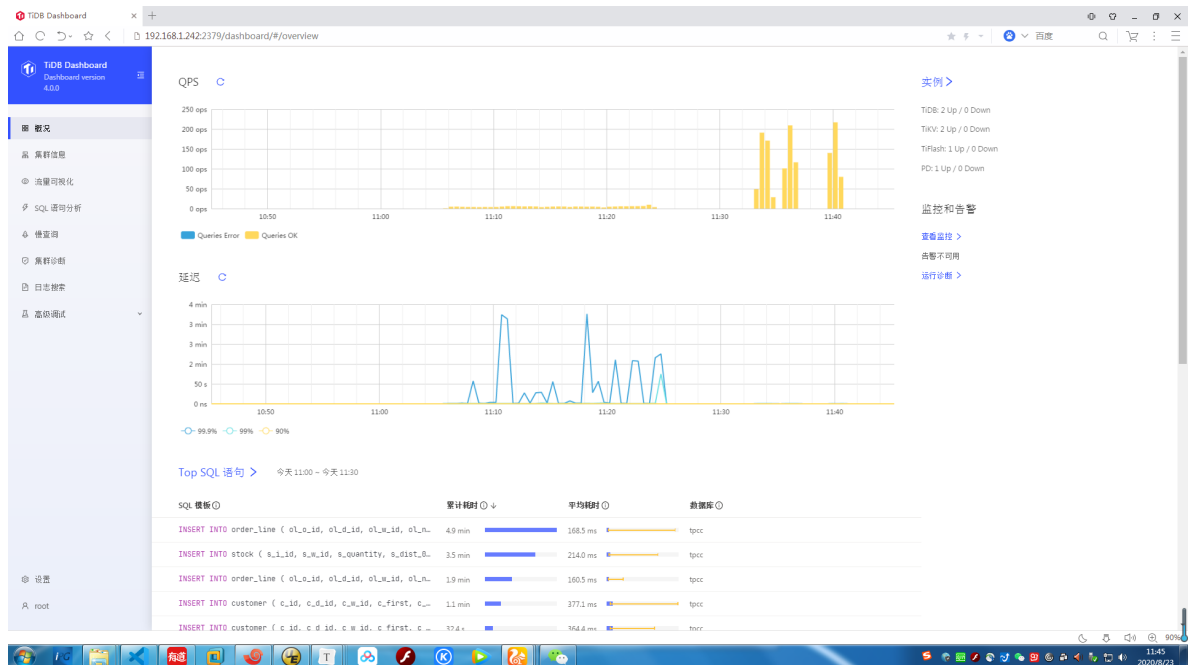
```
[root@tbase01 go-tpc]# ./bin/go-tpc tpcc -H 192.168.1.241 -P 4000 -D tpcc --
warehouses 10 run --time 1m --threads 5
[Current] DELIVERY - Takes(s): 8.6, Count: 7, TPM: 49.0, Sum(ms): 8367, Avg(ms):
1195, 90th(ms): 1500, 99th(ms): 1500, 99.9th(ms): 1500
[Current] NEW_ORDER - Takes(s): 9.5, Count: 62, TPM: 391.4, Sum(ms): 22975,
Avg(ms): 370, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] ORDER_STATUS - Takes(s): 8.8, Count: 7, TPM: 47.7, Sum(ms): 347,
Avg(ms): 49, 90th(ms): 160, 99th(ms): 160, 99.9th(ms): 160
[Current] PAYMENT - Takes(s): 9.7, Count: 51, TPM: 315.4, Sum(ms): 16324,
Avg(ms): 320, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] STOCK_LEVEL - Takes(s): 9.8, Count: 5, TPM: 30.7, Sum(ms): 400,
Avg(ms): 80, 90th(ms): 192, 99th(ms): 192, 99.9th(ms): 192
[Current] DELIVERY - Takes(s): 6.1, Count: 3, TPM: 29.3, Sum(ms): 3829, Avg(ms):
1276, 90th(ms): 1500, 99th(ms): 1500, 99.9th(ms): 1500
[Current] NEW_ORDER - Takes(s): 9.9, Count: 69, TPM: 416.4, Sum(ms): 25715,
Avg(ms): 372, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] ORDER_STATUS - Takes(s): 9.8, Count: 4, TPM: 24.5, Sum(ms): 110,
Avg(ms): 27, 90th(ms): 80, 99th(ms): 80, 99.9th(ms): 80
[Current] PAYMENT - Takes(s): 9.9, Count: 61, TPM: 369.1, Sum(ms): 19871,
Avg(ms): 325, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] STOCK_LEVEL - Takes(s): 7.8, Count: 4, TPM: 30.9, Sum(ms): 110,
Avg(ms): 27, 90th(ms): 48, 99th(ms): 48, 99.9th(ms): 48
[Current] DELIVERY - Takes(s): 9.7, Count: 7, TPM: 43.1, Sum(ms): 8333, Avg(ms):
1190, 90th(ms): 1500, 99th(ms): 1500, 99.9th(ms): 1500
[Current] NEW_ORDER - Takes(s): 9.9, Count: 71, TPM: 431.4, Sum(ms): 23857,
Avg(ms): 336, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] ORDER_STATUS - Takes(s): 9.9, Count: 9, TPM: 54.3, Sum(ms): 129,
Avg(ms): 14, 90th(ms): 80, 99th(ms): 80, 99.9th(ms): 80
[Current] PAYMENT - Takes(s): 9.9, Count: 65, TPM: 392.1, Sum(ms): 18103,
Avg(ms): 278, 90th(ms): 512, 99th(ms): 512, 99.9th(ms): 512
```

```

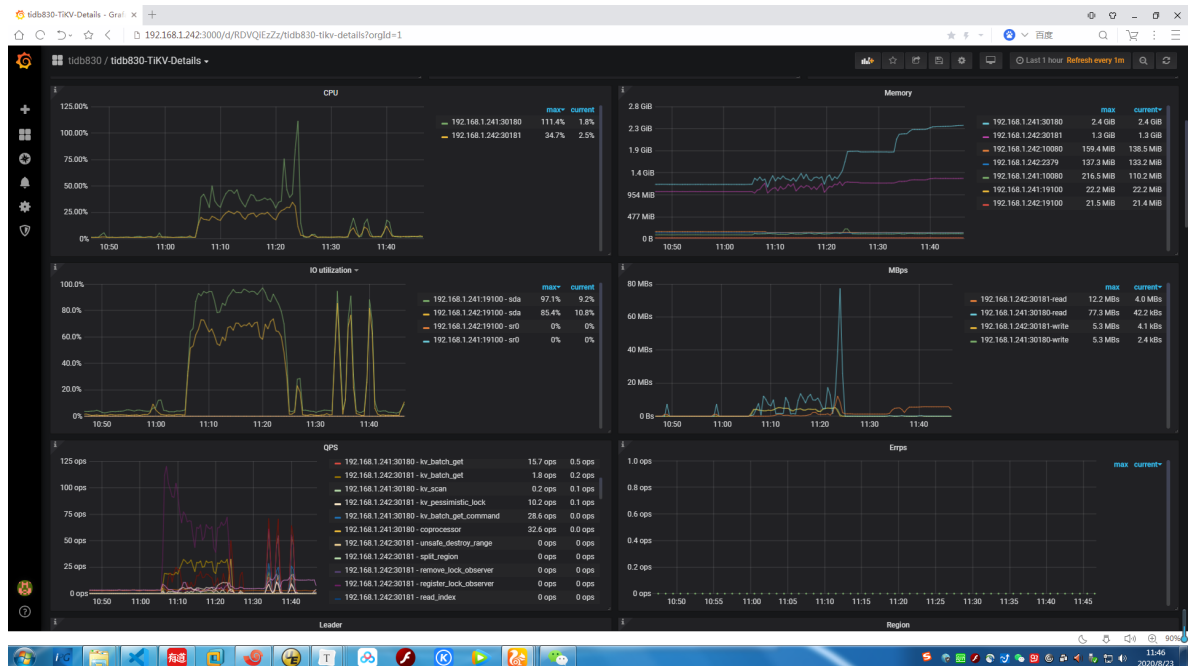
[Current] STOCK_LEVEL - Takes(s): 8.7, Count: 6, TPM: 41.6, Sum(ms): 90,
Avg(ms): 15, 90th(ms): 24, 99th(ms): 24, 99.9th(ms): 24
[Current] DELIVERY - Takes(s): 5.8, Count: 7, TPM: 73.0, Sum(ms): 8626, Avg(ms):
1232, 90th(ms): 4000, 99th(ms): 4000, 99.9th(ms): 4000
[Current] NEW_ORDER - Takes(s): 9.9, Count: 74, TPM: 449.9, Sum(ms): 26388,
Avg(ms): 356, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] ORDER_STATUS - Takes(s): 9.9, Count: 5, TPM: 30.4, Sum(ms): 39,
Avg(ms): 7, 90th(ms): 9, 99th(ms): 9, 99.9th(ms): 9
[Current] PAYMENT - Takes(s): 9.9, Count: 50, TPM: 303.0, Sum(ms): 14467,
Avg(ms): 289, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] STOCK_LEVEL - Takes(s): 9.3, Count: 6, TPM: 38.8, Sum(ms): 119,
Avg(ms): 19, 90th(ms): 48, 99th(ms): 48, 99.9th(ms): 48
[Current] DELIVERY - Takes(s): 7.4, Count: 7, TPM: 57.0, Sum(ms): 9067, Avg(ms):
1295, 90th(ms): 1500, 99th(ms): 1500, 99.9th(ms): 1500
[Current] NEW_ORDER - Takes(s): 9.9, Count: 64, TPM: 387.1, Sum(ms): 23306,
Avg(ms): 364, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] ORDER_STATUS - Takes(s): 9.9, Count: 6, TPM: 36.3, Sum(ms): 46,
Avg(ms): 7, 90th(ms): 9, 99th(ms): 9, 99.9th(ms): 9
[Current] PAYMENT - Takes(s): 9.7, Count: 58, TPM: 357.1, Sum(ms): 17969,
Avg(ms): 309, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Current] STOCK_LEVEL - Takes(s): 8.8, Count: 7, TPM: 47.9, Sum(ms): 95,
Avg(ms): 13, 90th(ms): 16, 99th(ms): 16, 99.9th(ms): 16
Finished
[Summary] DELIVERY - Takes(s): 58.6, Count: 37, TPM: 37.9, Sum(ms): 45744,
Avg(ms): 1236, 90th(ms): 1500, 99th(ms): 4000, 99.9th(ms): 4000
[Summary] NEW_ORDER - Takes(s): 59.5, Count: 409, TPM: 412.2, Sum(ms): 147470,
Avg(ms): 360, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Summary] NEW_ORDER_ERR - Takes(s): 59.5, Count: 3, TPM: 3.0, Sum(ms): 554,
Avg(ms): 184, 90th(ms): 256, 99th(ms): 256, 99.9th(ms): 256
[Summary] ORDER_STATUS - Takes(s): 58.8, Count: 39, TPM: 39.8, Sum(ms): 744,
Avg(ms): 19, 90th(ms): 80, 99th(ms): 160, 99.9th(ms): 160
[Summary] PAYMENT - Takes(s): 59.7, Count: 336, TPM: 337.5, Sum(ms): 103892,
Avg(ms): 309, 90th(ms): 512, 99th(ms): 1000, 99.9th(ms): 1000
[Summary] PAYMENT_ERR - Takes(s): 59.7, Count: 1, TPM: 1.0, Sum(ms): 43,
Avg(ms): 43, 90th(ms): 48, 99th(ms): 48, 99.9th(ms): 48
[Summary] STOCK_LEVEL - Takes(s): 59.8, Count: 37, TPM: 37.1, Sum(ms): 936,
Avg(ms): 25, 90th(ms): 48, 99th(ms): 192, 99.9th(ms): 192
tpmC: 412.2
[root@tbase01 go-tpc]#

```

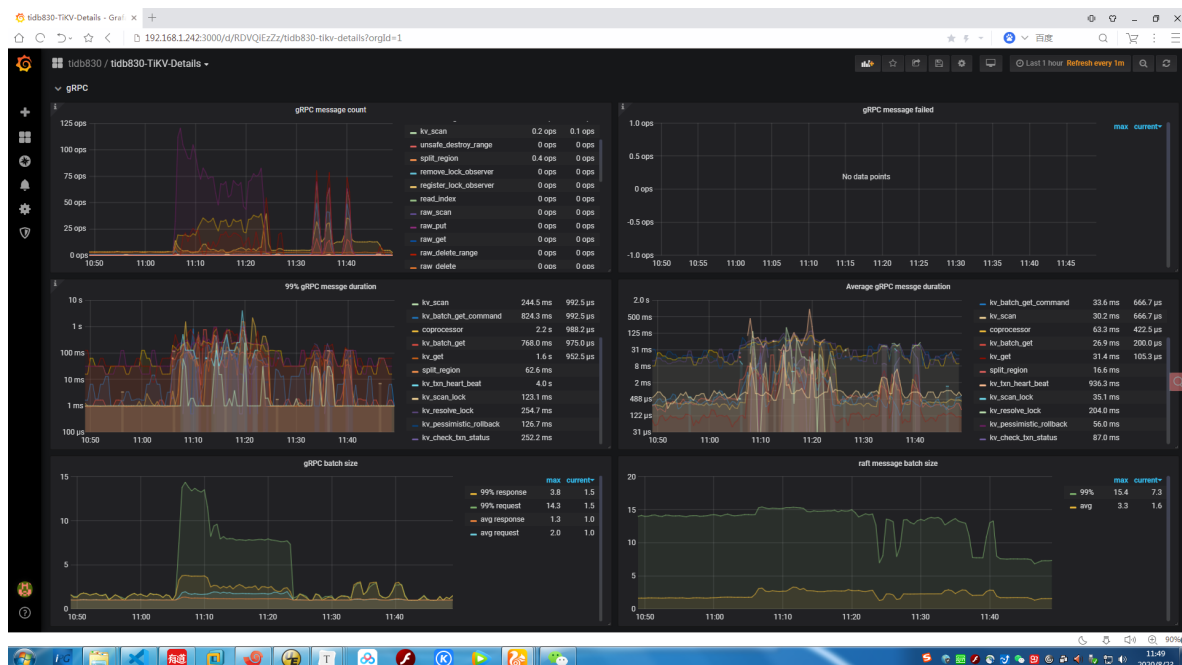
qps在200左右，延时在200ms左右。



集群server的状态，cpu在25%，内存缓慢上升,disk繁忙98%



集群gprc的状态:qps和是300左右（个体60），个体延时是100ms左右，平均延时30ms左右。压测时批处理gprc每组3gprc，批处理raft消息每批13个。



接近3M/s的磁盘写已经达到磁盘随机同步io上限：

```
Device:          rrqm/s  wrqm/s    r/s     w/s      rkB/s    wkB/s  avgrq-sz
avgqu-sz        await  r_await  w_await  svctm    %util
sda              0.00    140.00    0.00    305.00    0.00    2596.00  17.02
0.94    3.07    0.00    3.07    2.87    87.50
```

```
Device:          rrqm/s  wrqm/s    r/s     w/s      rkB/s    wkB/s  avgrq-sz
avgqu-sz        await  r_await  w_await  svctm    %util
sda              0.00    115.00    0.00    340.00    0.00    2648.00  15.58
0.97    2.81    0.00    2.81    2.72    92.60
```

```
Device:          rrqm/s  wrqm/s    r/s     w/s      rkB/s    wkB/s  avgrq-sz
avgqu-sz        await  r_await  w_await  svctm    %util
sda              0.00    117.00    0.00    302.00    0.00    1884.00  12.48
0.97    3.25    0.00    3.25    3.22    97.10
```


```
Device:          rrqm/s  wrqm/s    r/s     w/s      rkB/s    wkB/s  avgrq-sz
avgqu-sz        await  r_await  w_await  svctm    %util
sda              0.00    105.00    0.00    291.00    0.00    2820.00  19.38
0.94    3.22    0.00    3.22    3.15    91.60
```

```
^C
[root@tbase01 wangwei]#
```

写wal\_file用掉了300kB/s，占用1/8的磁盘性能消耗



## 磁盘



## RocksDB



## 优化方向

在保证“事务数据安全落盘”和尽量保证cap的前提下优化落盘机制。先记日志和分组打包是很好的方向。